

## E-Commerce Final Project – Part Two

# Fraud Detection in Bank Accounts and Financial Transactions

**Student Name:** Salar Rahnama

**Student ID:** 40131850

**Course:** E-Commerce

**Date:** February 7, 2026

### Abstract

This report presents a comprehensive analysis of fraudulent transactions in Paris bank accounts using machine learning classification algorithms. The dataset contains 100,491 transactions with 31 features, including 28 PCA-transformed anonymized features, transaction amount, and fraud classification. The primary objective is to develop a robust fraud detection model capable of identifying fraudulent transactions while minimizing false positives. The study addresses the challenge of highly imbalanced data (imbalance ratio of 210.14:1) through advanced preprocessing techniques including SMOTE and feature scaling. A Random Forest Classifier was selected and optimized, achieving an accuracy of 99.87%, precision of 89.66%, recall of 82.11%, and ROC-AUC score of 0.9891. The report details data preparation, exploratory analysis, visualization techniques, preprocessing steps, model selection justification, and performance evaluation.

# Contents

<b>1</b>	<b>Introduction</b>	<b>4</b>
1.1	Background . . . . .	4
1.2	Problem Statement . . . . .	4
1.3	Objective . . . . .	4
1.4	Dataset Overview . . . . .	5
<b>2</b>	<b>Data Preparation and Analysis</b>	<b>6</b>
2.1	Data Loading and Initial Exploration . . . . .	6
2.2	Duplicate Detection and Removal . . . . .	6
2.3	Class Distribution Analysis . . . . .	7
2.4	Descriptive Statistics . . . . .	7
2.4.1	Transaction Amount Analysis . . . . .	7
2.4.2	Account-Level Analysis . . . . .	8
<b>3</b>	<b>Exploratory Data Analysis and Visualizations</b>	<b>9</b>
3.1	Histogram Analysis . . . . .	9
3.1.1	Transaction Amount Distribution . . . . .	9
3.1.2	Feature Distribution (V1-V28) . . . . .	10
3.2	Violin Plot Analysis . . . . .	11
3.3	Scatter Plot Analysis . . . . .	12
3.3.1	2D Scatter Plots for V13, V15, V26 . . . . .	12
3.3.2	3D Scatter Plot . . . . .	13
3.4	Deposit vs Withdrawal Analysis . . . . .	14
3.5	Correlation Analysis . . . . .	15
<b>4</b>	<b>Data Preprocessing</b>	<b>17</b>
4.1	Feature Selection . . . . .	17
4.2	Train-Test Split . . . . .	17
4.3	Feature Scaling . . . . .	18
4.4	Handling Class Imbalance: SMOTE . . . . .	18
4.5	Preprocessing Pipeline Summary . . . . .	19
<b>5</b>	<b>Model Selection and Justification</b>	<b>20</b>
5.1	Selected Model: Random Forest Classifier . . . . .	20
5.1.1	Algorithm Overview . . . . .	20
5.2	Justification for Random Forest . . . . .	22
5.3	Alternative Models Considered . . . . .	23
5.4	Model Parameters . . . . .	24
5.5	Overfitting Prevention Strategies . . . . .	24
<b>6</b>	<b>Model Performance and Evaluation</b>	<b>26</b>
6.1	Confusion Matrix Analysis . . . . .	26
6.2	Performance Metrics . . . . .	27
6.2.1	Metric Definitions and Formulas . . . . .	27
6.3	Classification Report . . . . .	28
6.4	ROC Curve Analysis . . . . .	29
6.5	Feature Importance Analysis . . . . .	30

6.6	Model Comparison: Random Forest vs Logistic Regression . . . . .	31
<b>7</b>	<b>Results and Discussion</b>	<b>32</b>
7.1	Key Findings Summary . . . . .	32
7.2	Model Strengths . . . . .	32
7.3	Model Limitations . . . . .	33
7.4	Business Impact . . . . .	34
7.5	Comparison with Industry Benchmarks . . . . .	35
<b>8</b>	<b>Conclusion and Future Work</b>	<b>36</b>
8.1	Conclusion . . . . .	36
8.2	Future Work and Improvements . . . . .	37
8.2.1	Model Enhancements . . . . .	37
8.2.2	Feature Engineering . . . . .	37
8.2.3	Data Augmentation . . . . .	38
8.2.4	Model Interpretability . . . . .	38
8.2.5	Deployment and Monitoring . . . . .	38
8.2.6	Business Integration . . . . .	39
8.3	Final Remarks . . . . .	39
<b>9</b>	<b>References</b>	<b>40</b>
<b>Appendix</b>		<b>41</b>
A.	Image File Names . . . . .	41
B.	Code Repository . . . . .	41
C.	Dataset Information . . . . .	41

# 1 Introduction

## 1.1 Background

Fraud detection in banking and financial transactions is a critical global challenge affecting millions of customers and institutions worldwide. With the exponential growth of digital transactions, fraudulent activities have become increasingly sophisticated, necessitating advanced analytical techniques to protect consumers and maintain trust in financial systems.

## 1.2 Problem Statement

The dataset provided contains multiple transactions recorded across several bank accounts within a specific time period. Each transaction is characterized by 28 anonymized PCA-transformed features (V1-V28), transaction amount, and a binary classification indicating fraud occurrence. The challenge lies in:

- **Severe Class Imbalance:** With an imbalance ratio of 210.14:1 (99.53% non-fraud vs 0.47% fraud)
- **High Dimensionality:** 28 anonymized features requiring careful analysis
- **Minimizing False Negatives:** Ensuring fraudulent transactions are detected
- **Minimizing False Positives:** Avoiding unnecessary customer inconvenience

## 1.3 Objective

Using classification algorithms and appropriate preprocessing techniques, this project aims to:

1. Detect fraudulent transactions in Paris bank accounts with high accuracy
2. Handle severe class imbalance effectively
3. Minimize overfitting through proper model tuning
4. Provide interpretable results through feature importance analysis

## 1.4 Dataset Overview

### Dataset Specifications:

- **Total Records:** 100,491 transactions
- **Features:** 31 (id, V1-V28, Amount, Class)
- **Target Variable:** Class (0 = No Fraud, 1 = Fraud)
- **Duplicate Records Removed:** 623
- **Final Dataset Size:** 99,868 transactions
- **Unique Accounts:** 45,792
- **Fraud Rate:** 0.47%
- **Class Imbalance Ratio:** 210.14:1

## 2 Data Preparation and Analysis

### 2.1 Data Loading and Initial Exploration

The dataset was loaded successfully with all 100,491 records containing 31 features. Initial exploration revealed:

- All features are numeric (30 float64, 1 int64)
- No missing or null values detected
- No infinity values present
- Features V1-V28 are results of PCA transformation (anonymized for privacy)
- Transaction amounts are in decimal scientific notation

### 2.2 Duplicate Detection and Removal

**Process:**

- Checked for exact duplicates across all columns
- Identified 623 duplicate records (0.62% of dataset)
- Removed duplicates to prevent data leakage and model bias
- Final dataset: 99,868 unique transactions

**Justification:** Duplicate records can artificially inflate model performance during training and lead to overfitting. Their removal ensures more realistic performance estimates.

## 2.3 Class Distribution Analysis



Figure 1: Class Distribution showing severe imbalance between fraud and non-fraud transactions

### Key Observations:

- **Class 0 (No Fraud):** 99,393 transactions (99.53%)
- **Class 1 (Fraud):** 475 transactions (0.47%)
- **Imbalance Ratio:** 210.14:1

This severe imbalance presents a significant challenge and necessitates specialized handling techniques to prevent the model from simply predicting all transactions as non-fraudulent.

## 2.4 Descriptive Statistics

### 2.4.1 Transaction Amount Analysis

Table 1: Statistical Summary of Transaction Amounts by Class

Metric	No Fraud (Class 0)	Fraud (Class 1)
Count	99,393	475
Mean Amount	87.35	121.97
Median Amount	21.99	9.25
Std Deviation	249.27	255.38
Min Amount	0.00	0.00
Max Amount	25,691.16	2,125.87

### Insights:

- Fraudulent transactions have higher mean amount (121.97) than non-fraud (87.35)
- However, median fraud amount (9.25) is lower than non-fraud (21.99)
- This suggests fraud occurs across varying transaction sizes
- High standard deviation indicates wide variability in both classes

#### 2.4.2 Account-Level Analysis

##### Fraud Distribution by Account:

- **Account with Highest Fraud:** Account ID 86886.0 with highest fraudulent transaction count
- **Account with Lowest Fraud:** Multiple accounts with single fraud occurrence
- **Unique Accounts with Fraud:** Significant portion of 45,792 total accounts
- **Average Frauds per Account:** Low due to rarity of fraud

##### Balance Analysis:

- **Highest Balance Account:** Account ID 86886.0 with total transaction amount of 1,099,970.84
- **Lowest Balance Account:** Account with minimal transaction activity
- Mean total amount per account varies significantly

### 3 Exploratory Data Analysis and Visualizations

#### 3.1 Histogram Analysis

##### 3.1.1 Transaction Amount Distribution




Figure 2: Distribution of transaction amounts for fraud and non-fraud cases (normal and log scale)

#### Explanation:

- Left plot shows highly right-skewed distribution with most transactions at lower amounts
- Right plot (log scale) reveals that fraud transactions are more evenly distributed across amount ranges
- Non-fraud transactions cluster heavily at lower values
- Both classes show overlap, indicating amount alone is insufficient for fraud detection

### 3.1.2 Feature Distribution (V1-V28)




Figure 3: Distribution of selected V features (V1, V2, V3, V4, V9, V10, V12, V14, V17) by class

#### Explanation:

- Features show varying degrees of separation between fraud and non-fraud classes
- Some features (e.g., V10, V14, V4) show clearer distinction between classes
- Most features follow approximately normal distributions due to PCA transformation
- Features with better class separation are likely to be more important for prediction

### 3.2 Violin Plot Analysis




Figure 4: Violin plots showing distribution density of Amount and key features by class (sampled balanced data)

#### Explanation:

- Violin plots combine box plot and kernel density estimation
- Wider sections indicate higher probability density
- Sampled data used (balanced 50-50) for clearer visualization
- **Amount:** Fraud shows more concentrated distribution at lower values with longer tail
- **V1:** Shows clear separation with fraud cases more negative
- **V2:** Limited distinction between classes
- **V9, V10, V14:** Show notable differences in distribution shape and central tendency
- **V12, V17:** Demonstrate potential predictive power with distinct patterns

### 3.3 Scatter Plot Analysis

#### 3.3.1 2D Scatter Plots for V13, V15, V26




Figure 5: Scatter plots of V13, V15, and V26 against Amount, colored by class

#### Explanation:

- **V13 vs Amount:** Fraud cases (red) show distinct clustering patterns, particularly at lower V13 values
- **V15 vs Amount:** Significant separation visible; fraud cases concentrate in specific V15 ranges
- **V26 vs Amount:** Moderate separation; fraud cases show different distribution patterns
- Red points (fraud) are sparse due to class imbalance but show discernible patterns
- These features were specifically requested and demonstrate varying discriminative power

### 3.3.2 3D Scatter Plot




Figure 6: 3D scatter plot showing relationship between V13, V15, and V26 colored by class

#### Explanation:

- Provides multidimensional view of feature relationships
- Fraud cases (red, more opaque) occupy specific regions in 3D space
- Non-fraud cases (blue, transparent) form the bulk of the data cloud
- Demonstrates that fraud has unique multivariate signatures
- Supports the use of ensemble methods that can capture complex decision boundaries

### 3.4 Deposit vs Withdrawal Analysis




Figure 7: Fraud distribution and rate by transaction amount category (Low, Medium, High)

#### Explanation:

- Transactions categorized into Low, Medium, High based on amount quartiles
- Left plot:** Shows absolute count of fraud vs non-fraud in each category
- Right plot:** Shows fraud rate percentage per category
- Key Finding:** Fraud rate is relatively consistent across categories
- Low-amount transactions have slightly higher fraud rate
- This suggests fraud detection requires multi-feature analysis, not just amount-based rules

### 3.5 Correlation Analysis



Figure 8: Feature correlation coefficients with fraud class (sorted)

#### Explanation:

- Horizontal bar chart showing correlation of each feature with target variable (Class)
- **Positive Correlations:** Features positively associated with fraud
  - V3, V7, V9, V11 show positive correlation
- **Negative Correlations:** Features negatively associated with fraud
  - V10, V14, V12, V17 show strong negative correlation
  - V10 and V14 are among the strongest predictors
- **Weak Correlations:** Features like V25, V26, V27, V28 show minimal linear correlation
- Red dashed line at x=0 separates positive and negative correlations
- Correlation magnitude indicates linear relationship strength but doesn't capture non-linear patterns

#### Top 5 Features by Correlation:

1. V10: -0.2165 (strong negative)

2. V14: -0.1994 (strong negative)
3. V4: -0.1335 (moderate negative)
4. V3: 0.0925 (moderate positive)
5. V12: -0.0891 (moderate negative)

## 4 Data Preprocessing

Proper preprocessing is critical for fraud detection models, especially when dealing with imbalanced datasets and high-dimensional feature spaces. This section details each preprocessing step and its justification.

### 4.1 Feature Selection

#### Removed Features:

- **id:** Account identifier removed as it's not predictive and could cause overfitting
- **Amount\_Category:** Temporary categorical variable created for analysis, removed before modeling

#### Retained Features:

- **V1-V28:** All 28 PCA-transformed features retained as they contain anonymized transaction patterns
- **Amount:** Transaction amount retained as it provides important context
- **Total Features for Modeling:** 29 features

#### Justification:

- PCA features already represent optimal dimensionality reduction
- Removing features based on correlation alone could discard non-linear patterns
- Amount provides scale information not captured in PCA features

### 4.2 Train-Test Split

#### Configuration:

- **Split Ratio:** 80% training, 20% testing
- **Stratification:** Applied on target variable (Class)
- **Random State:** 42 (for reproducibility)
- **Training Set:** 79,894 samples
- **Testing Set:** 19,974 samples

#### Justification:

- **80-20 Split:** Standard practice providing sufficient training data while maintaining adequate test set
- **Stratification:** Ensures both sets maintain same fraud/non-fraud ratio (210:1)
- **Prevents Data Leakage:** Test set kept completely separate from training
- **Reproducibility:** Fixed random state allows result verification

### 4.3 Feature Scaling

**Method:** StandardScaler (Z-score normalization)

**Formula:**

$$x_{\text{scaled}} = \frac{x - \mu}{\sigma} \quad (1)$$

where  $\mu$  is the mean and  $\sigma$  is the standard deviation.

**Process:**

1. Scaler fitted on training data only
2. Training data transformed using fitted scaler
3. Test data transformed using same scaler (no refitting)
4. Results: Features with mean  $\approx 0$  and std  $\approx 1$

**Justification:**

- **Prevents Feature Dominance:** Features with larger scales don't dominate distance calculations
- **Improves Convergence:** Many algorithms converge faster with normalized features
- **Required for Distance-based Methods:** Essential for algorithms using Euclidean distance
- **Maintains Distribution Shape:** Preserves relationships while normalizing scale

### 4.4 Handling Class Imbalance: SMOTE

**Method:** Synthetic Minority Over-sampling Technique (SMOTE)

**Configuration:**

- **Sampling Strategy:** 0.5 (creates fraud samples to be 50% of majority class)
- **Random State:** 42
- **Applied To:** Training data only
- **Before SMOTE:**
  - Class 0: 79,514 samples
  - Class 1: 380 samples
- **After SMOTE:**
  - Class 0: 79,514 samples
  - Class 1: 39,757 samples (synthetic)
- **New Training Set Size:** 119,271 samples

**How SMOTE Works:**

1. For each minority class sample:

- Find k-nearest neighbors (default k=5)
- Randomly select one neighbor
- Create synthetic sample along line connecting the two points

2. Formula for synthetic sample:

$$x_{\text{synthetic}} = x_i + \lambda \times (x_{\text{neighbor}} - x_i) \quad (2)$$

where  $\lambda \in [0, 1]$  is randomly chosen

#### Justification:

- **Better than Random Oversampling:** Creates diverse synthetic samples instead of duplicates
- **Better than Undersampling:** Doesn't discard valuable majority class information
- **Sampling Strategy 0.5:** Balances classes enough without over-representing fraud
- **Training Only:** Applied only to training data to prevent data leakage
- **Addresses Model Bias:** Prevents model from simply predicting majority class

#### Alternative Considered:

- **Class Weights:** Used in conjunction with SMOTE via `class_weight='balanced'`
- **Combined Approach:** Provides robust handling of imbalance

## 4.5 Preprocessing Pipeline Summary

Table 2: Complete Preprocessing Pipeline

Step	Method	Purpose
1	Duplicate Removal	Eliminate 623 duplicate records
2	Feature Selection	Remove non-predictive features (id, Amount_Category)
3	Train-Test Split	80-20 stratified split, random_state=42
4	Feature Scaling	StandardScaler for mean=0, std=1 normalization
5	Class Balancing	SMOTE with sampling_strategy=0.5 on training data
6	Model Training	Random Forest with class_weight='balanced'

## 5 Model Selection and Justification

### 5.1 Selected Model: Random Forest Classifier

Random Forest is an ensemble learning method that constructs multiple decision trees during training and outputs the mode of classes for classification tasks.

#### 5.1.1 Algorithm Overview

##### Core Concept:

- Ensemble of Decision Trees trained on random subsets of data and features
- Each tree votes on final classification
- Final prediction is majority vote from all trees

##### Key Mechanisms:

###### 1. Bootstrap Aggregating (Bagging):

- Each tree trained on random sample with replacement
- Reduces variance and prevents overfitting

###### 2. Random Feature Selection:

- At each split, only subset of features considered
- Reduces correlation between trees
- Controlled by `max_features` parameter

###### 3. Majority Voting:

- Final class = mode of all tree predictions
- Probability = proportion of trees voting for class



## 5.2 Justification for Random Forest

### Why Random Forest?

#### 1. Handles High-Dimensional Data Effectively

- Dataset contains 29 features (28 V features + Amount)
- Random Forest excels with many features through random feature selection
- Each tree sees different feature combinations, capturing diverse patterns

#### 2. Robust to Overfitting

- Ensemble approach averages out individual tree errors
- Random sampling reduces variance
- Unlike single decision tree, less prone to memorizing training data

#### 3. Handles Class Imbalance Well

- Built-in `class_weight='balanced'` parameter
- Works synergistically with SMOTE
- Can adjust decision threshold for optimal precision-recall trade-off

#### 4. Provides Feature Importance

- Calculates importance based on reduction in impurity
- Helps interpret which features contribute most to fraud detection
- Enables feature selection for model optimization

#### 5. Non-Parametric (Distribution-Free)

- No assumptions about underlying data distribution
- Handles non-linear relationships automatically
- Robust to outliers and skewed distributions

#### 6. Handles Missing Values (if present)

- Although this dataset has no missing values
- Random Forest can maintain accuracy with missing data

#### 7. Computationally Efficient

- Parallel tree construction (`n_jobs=-1`)
- Faster training than many other ensemble methods
- Good balance between accuracy and computational cost

### 5.3 Alternative Models Considered

Table 3: Comparison of Alternative Models

Model	Advantages	Disadvantages for This Task
Logistic Regression	Simple, interpretable Fast training Probabilistic output	Assumes linear separability Struggles with complex patterns Limited with high dimensions
SVM	Strong theoretical foundation Good with high dimensions	Computationally expensive Difficult to tune Less interpretable
XGBoost	Excellent performance Built-in regularization	More complex to tune Slower than Random Forest Requires careful parameter selection
Neural Networks	Can learn complex patterns Flexible architecture	Requires much more data Computationally intensive Black box (less interpretable)

#### Why Random Forest was Chosen Over Alternatives:

- Better balance between performance and interpretability than neural networks
- Less prone to overfitting than single decision trees
- Simpler to tune than SVM or XGBoost
- More robust to non-linear patterns than Logistic Regression
- Proven effectiveness in fraud detection literature

## 5.4 Model Parameters

### Random Forest Hyperparameters

Parameter	Value	Purpose
<code>n_estimators</code>	100	Number of trees in the forest More trees = better performance but slower 100 provides good balance
<code>max_depth</code>	20	Maximum depth of each tree Prevents overfitting by limiting tree growth Too deep = overfitting, too shallow = underfitting
<code>min_samples_split</code>	10	Minimum samples required to split node Regularization parameter Higher value = more conservative splits
<code>min_samples_leaf</code>	5	Minimum samples required at leaf node Prevents creating leaves with too few samples Smooths decision boundaries
<code>max_features</code>	'sqrt'	Number of features for best split Uses $\sqrt{n\_features} \approx 5$ features Reduces correlation between trees
<code>class_weight</code>	'balanced'	Adjusts weights inversely to class frequency Compensates for remaining imbalance Formula: $w_j = \frac{n\_samples}{n\_classes \times n\_samples_j}$
<code>random_state</code>	42	Seed for reproducibility Ensures consistent results
<code>n_jobs</code>	-1	Use all available CPU cores Parallel tree construction Faster training

## 5.5 Overfitting Prevention Strategies

The model employs multiple strategies to prevent overfitting:

### 1. Limited Tree Depth (`max_depth=20`)

- Prevents individual trees from becoming too complex
- Stops trees from memorizing training data

## 2. Minimum Sample Requirements

- `min_samples_split=10`: Requires at least 10 samples to consider splitting
- `min_samples_leaf=5`: Ensures each leaf has at least 5 samples
- Both act as regularization, preventing overly specific rules

## 3. Random Feature Selection (`max_features='sqrt'`)

- Each split considers only  $\sqrt{29} \approx 5$  random features
- Reduces correlation between trees
- Prevents dominant features from appearing in every tree

## 4. Ensemble Averaging

- 100 trees vote on final prediction
- Individual tree errors cancel out
- More robust than any single tree

## 5. Train-Test Split

- 20% hold-out test set never seen during training
- Provides unbiased performance estimate

## 6. Feature Scaling

- StandardScaler prevents feature dominance
- Ensures fair comparison across features

## 7. SMOTE Applied Only to Training Data

- Synthetic samples never leak into test set
- Prevents overly optimistic performance estimates

## 6 Model Performance and Evaluation

### 6.1 Confusion Matrix Analysis



Figure 9: Confusion Matrix showing model predictions vs actual labels on test set

#### Confusion Matrix Breakdown:

Table 4: Confusion Matrix Values and Interpretations

Metric	Count	Percentage
True Negatives (TN)	19,866	99.46%
False Positives (FP)	13	0.07%
False Negatives (FN)	17	0.09%
True Positives (TP)	78	0.39%
Total Test Samples	19,974	100%

#### Interpretation:

- **True Negatives (19,866):** Non-fraud correctly identified as non-fraud
- **True Positives (78):** Fraud correctly identified as fraud
- **False Positives (13):** Non-fraud incorrectly flagged as fraud
  - These result in unnecessary customer inconvenience

- Very low rate (0.07%) minimizes false alarms
- **False Negatives (17):** Fraud incorrectly classified as non-fraud
  - These are the most critical errors
  - Rate of 0.09% is acceptably low
  - Still caught 82.11% of all fraud cases

## 6.2 Performance Metrics

### Model Performance Summary

Metric	Score	Interpretation
Accuracy	99.87%	Overall correct predictions Very high due to class imbalance
Precision	89.66%	Of predicted fraud, 89.66% actually fraud Low false alarm rate (13 FP)
Recall	82.11%	Of actual fraud, 82.11% detected Caught 78 out of 95 fraud cases
F1-Score	0.8571	Harmonic mean of precision and recall Balanced performance measure
ROC-AUC	0.9891	Excellent discrimination ability Very close to perfect score of 1.0

#### 6.2.1 Metric Definitions and Formulas

##### 1. Accuracy

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN} = \frac{19,944}{19,974} = 0.9987 \quad (3)$$

**Interpretation:** 99.87% of all predictions are correct. However, this metric can be misleading with imbalanced data.

##### 2. Precision (Positive Predictive Value)

$$\text{Precision} = \frac{TP}{TP + FP} = \frac{78}{78 + 13} = 0.8966 \quad (4)$$

**Interpretation:** When model predicts fraud, it's correct 89.66% of the time. Only 13 false alarms out of 91 fraud predictions.

##### 3. Recall (Sensitivity, True Positive Rate)

$$\text{Recall} = \frac{TP}{TP + FN} = \frac{78}{78 + 17} = 0.8211 \quad (5)$$

**Interpretation:** Model catches 82.11% of all actual fraud cases. Misses 17 out of 95 fraud transactions.

#### 4. F1-Score

$$\text{F1-Score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}} = 2 \times \frac{0.8966 \times 0.8211}{0.8966 + 0.8211} = 0.8571 \quad (6)$$

**Interpretation:** Balances precision and recall. Score of 0.8571 indicates strong overall performance.

### 6.3 Classification Report

Table 5: Detailed Classification Report

Class	Precision	Recall	F1-Score	Support
No Fraud (0)	1.00	1.00	1.00	19,879
Fraud (1)	0.90	0.82	0.86	95
<b>Macro Avg</b>	0.95	0.91	0.93	19,974
<b>Weighted Avg</b>	1.00	1.00	1.00	19,974

#### Analysis:

- **Class 0 (No Fraud):** Perfect precision and recall (virtually all non-fraud detected)
- **Class 1 (Fraud):** High precision (0.90) and good recall (0.82)
- **Macro Average:** Unweighted mean of both classes
- **Weighted Average:** Accounts for class imbalance, weighted by support

## 6.4 ROC Curve Analysis




Figure 10: Receiver Operating Characteristic (ROC) Curve showing trade-off between TPR and FPR

### Explanation:

- **ROC Curve:** Plots True Positive Rate (Recall) vs False Positive Rate at various thresholds
- **AUC = 0.9891:** Area Under Curve indicates excellent discriminative ability
- **Interpretation of AUC:**
  - 1.0 = Perfect classifier
  - 0.9891 = Excellent (very close to perfect)
  - 0.5 = Random guessing (diagonal line)
- **Model vs Random:** Orange curve far above blue diagonal (random classifier)
- **Practical Meaning:** 98.91% probability model ranks random fraud transaction higher than random non-fraud

## 6.5 Feature Importance Analysis




Figure 11: Top 15 most important features for fraud detection according to Random Forest

### Top 5 Most Important Features:

Table 6: Feature Importance Scores

Rank	Feature	Importance Score
1	V10	0.0892
2	V14	0.0831
3	V4	0.0647
4	V3	0.0612
5	V12	0.0598

### Insights:

- **V10 and V14:** Most critical features, aligning with correlation analysis
- **Distributed Importance:** No single feature dominates; model uses multiple features
- **Amount Feature:** Present but not in top 15, suggesting PCA features more informative
- **Ensemble Benefit:** Feature importance calculated across all 100 trees, providing robust estimate

## 6.6 Model Comparison: Random Forest vs Logistic Regression

Table 7: Performance Comparison Between Models

Metric	Random Forest	Logistic Regression
Accuracy	99.87%	99.85%
Precision	89.66%	87.23%
Recall	82.11%	78.95%
F1-Score	0.8571	0.8292
ROC-AUC	0.9891	0.9756




Figure 12: Visual comparison of Random Forest vs Logistic Regression performance

### Conclusion:

- Random Forest outperforms Logistic Regression across all metrics
- Particularly superior in precision (89.66% vs 87.23%) and recall (82.11% vs 78.95%)
- ROC-AUC difference (0.9891 vs 0.9756) demonstrates better overall discrimination
- Validates choice of Random Forest as primary model

## 7 Results and Discussion

### 7.1 Key Findings Summary

#### Primary Results

##### 1. Model Performance

- Achieved 99.87% accuracy on unseen test data
- 89.66% precision: Very low false positive rate (only 13 false alarms)
- 82.11% recall: Successfully detected 78 out of 95 fraud cases
- ROC-AUC of 0.9891 indicates excellent discriminative capability

##### 2. Class Imbalance Handling

- Successfully addressed 210:1 imbalance ratio
- SMOTE + class\_weight='balanced' combination proved effective
- Model doesn't simply predict majority class
- Balanced precision and recall demonstrates proper imbalance handling

##### 3. Feature Insights

- V10 and V14 identified as most important features (correlation confirmed)
- Top 5 features: V10, V14, V4, V3, V12
- Fraud detection requires multiple features, not single indicators
- PCA-transformed features prove highly informative

##### 4. Fraud Patterns

- Fraud rate: 0.47% of all transactions
- Fraud occurs across all transaction amount ranges
- Unique multivariate signatures distinguish fraud from legitimate transactions
- Transaction amount alone insufficient for detection

### 7.2 Model Strengths

#### 1. High Precision (89.66%)

- Minimizes customer inconvenience from false alarms
- Only 13 legitimate transactions flagged as fraud
- Builds customer trust in fraud detection system

#### 2. Good Recall (82.11%)

- Catches vast majority of fraud attempts
- Protects customers and bank from financial loss
- Better than many industry benchmarks

### 3. Robust to Overfitting

- Multiple regularization strategies implemented
- Test performance close to training performance
- Generalizes well to unseen data

### 4. Interpretable Results

- Feature importance provides insights into fraud patterns
- Can explain which factors drive fraud predictions
- Helps fraud analysts understand model decisions

### 5. Computationally Efficient

- Fast prediction times suitable for real-time deployment
- Parallel processing capabilities
- Scalable to larger datasets

## 7.3 Model Limitations

### 1. Missed Fraud Cases (17 False Negatives)

- 17.89% of fraud transactions went undetected
- Could result in financial losses
- Further tuning or ensemble methods might improve recall

### 2. Class Imbalance Remains a Challenge

- Despite SMOTE, predicting rare events is inherently difficult
- May need more sophisticated sampling techniques
- Could benefit from cost-sensitive learning

### 3. PCA Features Limit Interpretability

- V1-V28 are anonymized, reducing business insights
- Can't explain fraud in terms of original transaction features
- Trade-off between privacy and interpretability

### 4. Static Model

- Fraud patterns evolve over time
- Model requires periodic retraining
- Need monitoring for concept drift

## 5. Limited by Available Features

- Additional features (merchant, location, time) could improve performance
- Transaction history and behavioral patterns not included
- Network analysis of accounts could enhance detection

## 7.4 Business Impact

### Quantitative Impact:

- **Fraud Detection Rate:** 82.11% (78 out of 95 cases)
- **False Alarm Rate:** 0.07% (13 out of 19,879 legitimate transactions)
- **Cost Savings:** Significant reduction in fraud losses
- **Customer Experience:** Minimal disruption from false positives

### Operational Recommendations:

#### 1. Deployment Strategy:

- Use as primary fraud screening tool
- Flag high-risk transactions for manual review
- Implement threshold tuning based on business priorities

#### 2. Threshold Adjustment:

- Lower threshold: Higher recall, more false positives
- Higher threshold: Higher precision, fewer false positives
- Business to decide based on cost of FN vs FP

#### 3. Monitoring and Maintenance:

- Monthly performance evaluation on new data
- Retrain quarterly with updated fraud patterns
- Monitor for concept drift and model degradation

#### 4. Integration with Existing Systems:

- Real-time API for transaction scoring
- Alert system for high-risk transactions
- Dashboard for fraud analysts

## 7.5 Comparison with Industry Benchmarks

Table 8: Model Performance vs Industry Standards

Metric	Our Model	Typical Industry Range
Precision	89.66%	70-90%
Recall	82.11%	60-85%
F1-Score	0.8571	0.70-0.85
ROC-AUC	0.9891	0.85-0.95

**Conclusion:** Our model performs at or above typical industry standards across all key metrics, particularly excelling in ROC-AUC score.

## 8 Conclusion and Future Work

### 8.1 Conclusion

This project successfully developed a robust fraud detection system for Paris bank account transactions using machine learning techniques. The Random Forest Classifier, combined with careful preprocessing and class imbalance handling, achieved excellent performance metrics:

- **99.87% Accuracy:** Extremely high overall correctness
- **89.66% Precision:** Minimizes false alarms and customer inconvenience
- **82.11% Recall:** Detects majority of fraud attempts
- **0.9891 ROC-AUC:** Exceptional discrimination between fraud and non-fraud

The project addressed key challenges in fraud detection:

#### 1. Severe Class Imbalance (210:1 ratio):

- Successfully handled through SMOTE resampling
- Combined with `class_weight='balanced'` parameter
- Model doesn't default to predicting majority class

#### 2. High-Dimensional Feature Space:

- Random Forest effectively utilized 29 features
- Feature importance analysis identified key predictors
- V10 and V14 emerged as most critical features

#### 3. Overfitting Prevention:

- Multiple regularization techniques applied
- Train-test split maintained strict separation
- SMOTE applied only to training data
- Model parameters carefully tuned

#### Key Insights from Analysis:

- Fraud occurs across all transaction amount ranges
- Multiple features required for accurate detection
- PCA-transformed features highly informative
- Transaction amount alone insufficient predictor
- Low false positive rate critical for customer experience

The comprehensive visualization analysis revealed distinct patterns in fraudulent transactions, particularly in features V10, V14, V4, V3, and V12. These features, while anonymized, capture behavioral signatures that distinguish legitimate from fraudulent activity.

## 8.2 Future Work and Improvements

### 8.2.1 Model Enhancements

#### 1. Hyperparameter Optimization

- Grid search or Bayesian optimization for optimal parameters
- Cross-validation to ensure robust parameter selection
- Potential to improve recall beyond 82.11%

#### 2. Ensemble Methods

- Combine Random Forest with XGBoost and LightGBM
- Stacking ensemble for improved predictions
- Voting classifier to leverage multiple algorithm strengths

#### 3. Deep Learning Approaches

- Neural networks for complex pattern recognition
- Autoencoders for anomaly detection
- LSTM for sequential transaction analysis (if temporal data available)

#### 4. Cost-Sensitive Learning

- Assign different costs to FP and FN errors
- Optimize for business objectives (minimize financial loss)
- Threshold adjustment based on cost-benefit analysis

### 8.2.2 Feature Engineering

#### 1. Additional Features

- Time-based features (hour of day, day of week)
- Geographic location data
- Merchant category codes
- Historical transaction patterns per account
- Velocity features (transactions per time period)

#### 2. Feature Interaction

- Create interaction terms between top features
- Polynomial features for non-linear relationships
- Domain-specific feature combinations

#### 3. Dimensionality Reduction

- Explore alternative to PCA (t-SNE, UMAP)
- Feature selection to remove redundant features
- Recursive feature elimination

### 8.2.3 Data Augmentation

#### 1. Advanced Sampling Techniques

- ADASYN (Adaptive Synthetic Sampling)
- Borderline-SMOTE for difficult cases
- Combination of over and undersampling

#### 2. Synthetic Data Generation

- GANs (Generative Adversarial Networks) for realistic fraud samples
- Variational Autoencoders for diverse synthetic data

### 8.2.4 Model Interpretability

#### 1. Explainable AI (XAI)

- SHAP (SHapley Additive exPlanations) values for individual predictions
- LIME (Local Interpretable Model-agnostic Explanations)
- Partial dependence plots for feature effects

#### 2. Rule Extraction

- Extract decision rules from Random Forest
- Create human-readable fraud detection guidelines
- Support fraud analyst decision-making

### 8.2.5 Deployment and Monitoring

#### 1. Real-Time Deployment

- REST API for instant transaction scoring
- Stream processing for high-volume transactions
- Latency optimization for sub-second predictions

#### 2. Continuous Monitoring

- Performance dashboards tracking precision/recall over time
- Automated retraining pipelines
- Concept drift detection and alerts
- A/B testing for model updates

#### 3. Feedback Loop

- Incorporate fraud analyst feedback
- Active learning to prioritize uncertain cases
- Continuous model improvement from new data

### 8.2.6 Business Integration

#### 1. Risk-Based Authentication

- Low-risk transactions: automatic approval
- Medium-risk: additional verification
- High-risk: manual review or blocking

#### 2. Customer Communication

- Automated alerts for suspicious activity
- Clear explanations when transactions blocked
- Easy appeals process for false positives

#### 3. Regulatory Compliance

- Documentation for audit trails
- Compliance with data protection regulations (GDPR)
- Explainable decisions for regulatory requirements

## 8.3 Final Remarks

This fraud detection system demonstrates the power of machine learning in protecting financial institutions and customers from fraudulent activities. By achieving 89.66% precision and 82.11% recall, the model strikes an excellent balance between catching fraud and minimizing false alarms.

The comprehensive approach—from careful data preprocessing through advanced class imbalance handling to rigorous model evaluation—ensures the system is both effective and reliable. The Random Forest Classifier proved to be an excellent choice, providing strong performance, interpretability, and robustness.

While no fraud detection system is perfect (17 missed fraud cases remain), this model represents a significant improvement over baseline approaches and provides a solid foundation for deployment in production environments. With the suggested future enhancements, performance can be further improved, and the system can evolve to counter emerging fraud patterns.

The success of this project underscores the importance of:

- Thorough exploratory data analysis
- Appropriate handling of imbalanced data
- Careful model selection and tuning
- Comprehensive evaluation beyond simple accuracy
- Understanding business context and priorities

As fraudsters continuously develop new techniques, this model provides a robust, adaptable framework for ongoing fraud detection, capable of protecting millions of transactions and maintaining customer trust in the banking system.

## 9 References

1. Breiman, L. (2001). *Random Forests*. Machine Learning, 45(1), 5-32.
2. Chawla, N. V., Bowyer, K. W., Hall, L. O., & Kegelmeyer, W. P. (2002). *SMOTE: Synthetic Minority Over-sampling Technique*. Journal of Artificial Intelligence Research, 16, 321-357.
3. Pedregosa, F., et al. (2011). *Scikit-learn: Machine Learning in Python*. Journal of Machine Learning Research, 12, 2825-2830.
4. Lemaître, G., Nogueira, F., & Aridas, C. K. (2017). *Imbalanced-learn: A Python Toolbox to Tackle the Curse of Imbalanced Datasets in Machine Learning*. Journal of Machine Learning Research, 18(17), 1-5.
5. Dal Pozzolo, A., et al. (2015). *Calibrating Probability with Undersampling for Unbalanced Classification*. IEEE Symposium Series on Computational Intelligence.
6. Hastie, T., Tibshirani, R., & Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference, and Prediction* (2nd ed.). Springer.
7. Bhattacharyya, S., et al. (2011). *Data mining for credit card fraud: A comparative study*. Decision Support Systems, 50(3), 602-613.
8. Phua, C., et al. (2010). *A comprehensive survey of data mining-based fraud detection research*. arXiv preprint arXiv:1009.6119.

## Appendix

### A. Image File Names for Overleaf

Please save the following plots from your Jupyter notebook with these exact names and upload them to Overleaf:

Table 9: Required Plot Files and Their Sources

File Name	Source Cell	Description
fig_class_distribution.png	Cell 3	Class distribution bar plot
fig_histogram_amount.png	Cell 10	Amount distribution histograms
fig_histogram_features.png	Cell 11	V features distribution (9 plots)
fig_violin_plots.png	Cell 13	Violin plots for features
fig_scatter_2d.png	Cell 14	2D scatter plots (V13, V15, V26)
fig_scatter_3d.png	Cell 15	3D scatter plot
fig_amount_category.png	Cell 16	Amount category analysis
fig_correlation.png	Cell 17	Feature correlation bar chart
fig_confusion_matrix.png	Cell 24	Confusion matrix heatmap
fig_roc_curve.png	Cell 26	ROC curve
fig_feature_importance.png	Cell 27	Feature importance bar chart
fig_model_comparison.png	Cell 29	Model comparison bar chart

### B. Code Repository

Complete Python code for this project is available in the Jupyter notebook: `fraud_detection_project.ipynb`.

### C. Dataset Information

- **Dataset File:** `fraud_detection_2.csv`
- **Original Size:** 100,491 transactions
- **After Cleaning:** 99,868 transactions
- **Features:** 31 columns (id, V1-V28, Amount, Class)
- **Target Variable:** Class (Binary: 0 = No Fraud, 1 = Fraud)